

B-Tree Index Bloat

2019-10-03 | Jeroen Soesbergen

Let's introduce myself...

- Jeroen Soesbergen
- Scrum Master / Software Engineer
- Regular guest at the PostgreSQL User Group NL
- Learned a lot here
- Really nice to contribute back today



About Simacan

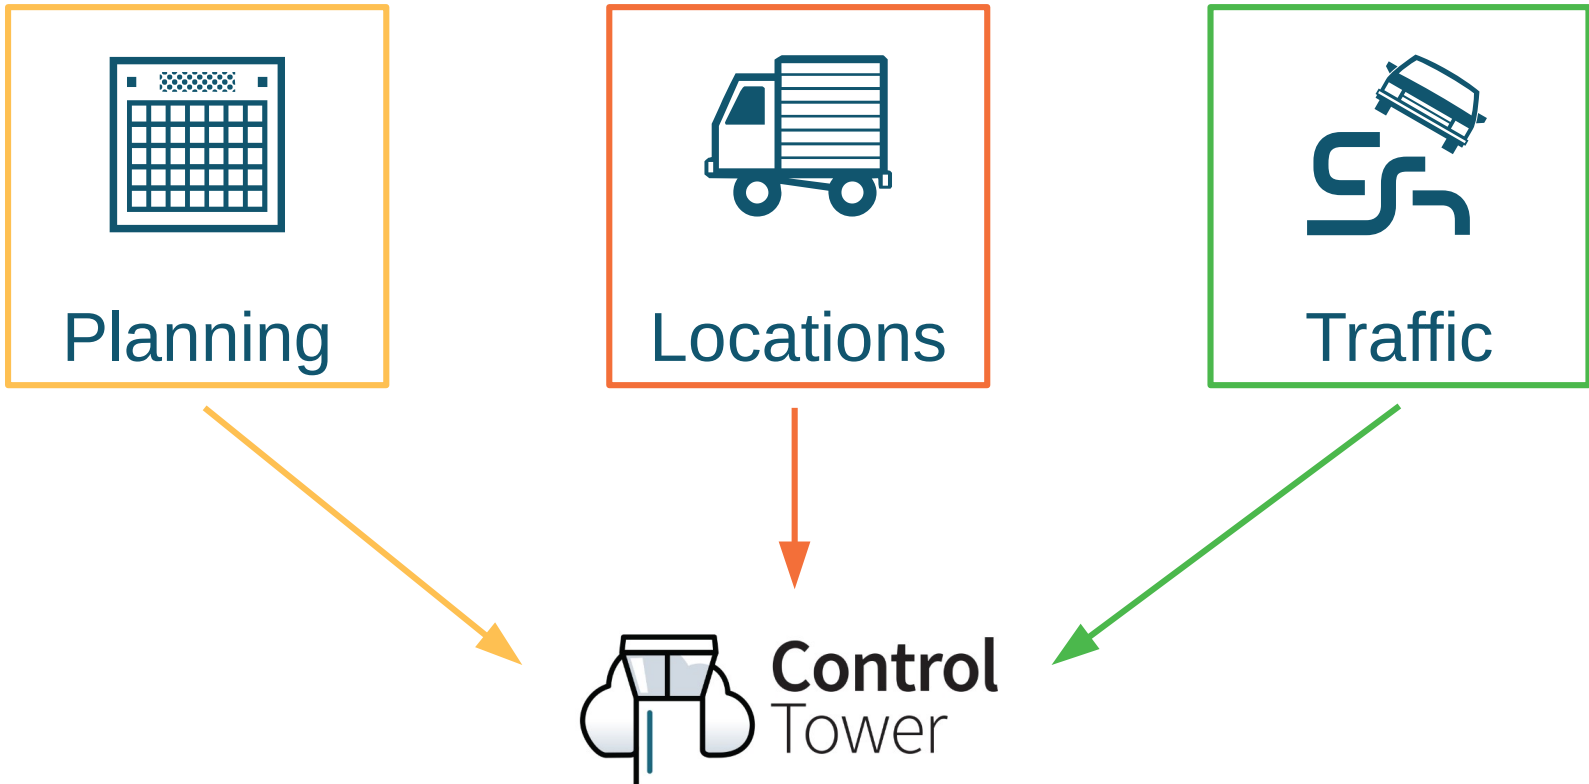
- Software for Traffic & Transport
- Our product:



Control
Tower



Simacan Control Tower (1)



Simacan Control Tower (2)



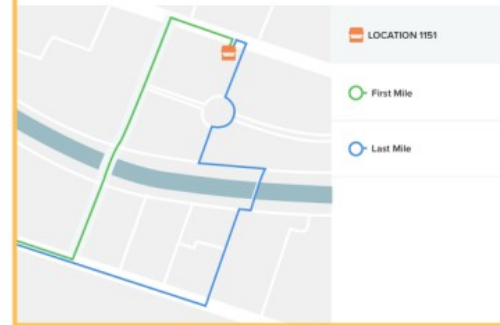
Customer Support, Planning, Operations



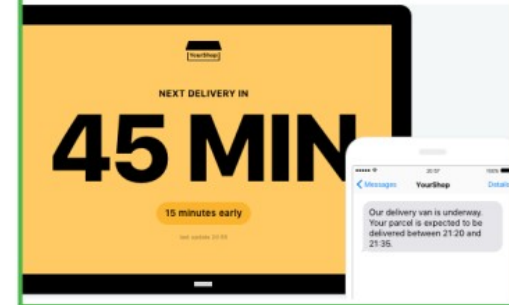
Distribution Centers



Carriers



Clients



Technology @ Simacan

- Microservice Cluster in AWS
- Relational database: PostgreSQL
 - Open Source
 - Spatial extension: PostGIS
 - Proven to be a reliable choice



B-Tree Index Bloat

- What is index bloat?
- Demo
- Where does it come from?
- How to deal with index bloat?



What is B-Tree Index Bloat?

- What is this, exactly?
 - B-Tree Index = default type of index
 - Bloat = an object on disk contains empty space
- Why is this a problem?
 - Waste of disk space and I/O
 - Performance impact
- How bad can it get?



Demo: How bad can it get?

- Take a simple table/index
- Perform update/delete/insert/vacuum
 - In a specific pattern
 - Exact pattern will be revealed later!
- Monitor table and index bloat

https://github.com/pgexperts/pgx_scripts



Where does index bloat come from?

- Index fragmentation
- Multi Version Concurrency Control (MVCC)
- Vacuum characteristics



Index fragmentation

- Indexes consist of 8kB pages
- When an index page is full:
 - It will be split
- Pages will never be merged
- Lower fill factor can prevent full pages



Multi Version Concurrency Control

- After a tuple is updated:
 - The old version still exists
 - A copy is created with the new values
- Both copies must be in the index!
- The index grows...



Vacuum characteristics

- VACUUM: Cleans up dead tuples
- Makes space in tables and index pages available for re-use
- Autovacuum daemon:
 - Triggered by vacuum threshold
 - Percentage of dead tuples in table
 - Works fine for tables, not always for indexes
 - Table tuples do not have to be ordered, but indexes do
 - Can be tweaked, when necessary

```
vacuum threshold = vacuum base threshold +  
                    vacuum scale factor * number of tuples
```



Demo: Revisit

- Let's check the demo!



Demo: Observations

- Very little table bloat
- The index goes crazy
- Interesting:
 - Index bloat can exist without table bloat
 - You need to monitor both separately



Demo: How does it work?

- Bloat 1% of the id range at a time
 - First: 10 updates on every tuple
 - Autovacuum not triggered by default
 - All copies must be indexed => Index bloat
 - Delete 80% of tuples
 - Bloats index even more
 - Vacuum
 - To 'clear' the updates and deletes
 - Insert new tuples, to fill table again
- This 1% of the index gets bloated badly
- Repeat on the next 1%



How to deal with index bloat?

- Prevent:
 - Tweak index fillfactor
 - Tweak autovacuum
- Fix:
 - REINDEX
 - Create replacement index



REINDEX

- Pretty simple
- Index contents are rebuilt from scratch
- Replaces the old one
- Table lock: SHARE
 - No inserts/updates/deletes
- Index lock: ACCESS EXCLUSIVE
 - Blocks all access (also reads) to the index
- This probably means downtime...



Create a replacement index

- Create an exact copy of the bloated index
- CREATE INDEX CONCURRENTLY
- Swap bloated and new index
- Extension: pg_repack
 - Does something like this



CREATE INDEX CONCURRENTLY

- Table lock: SHARE UPDATE EXCLUSIVE
 - No schema changes
 - No vacuum runs
 - No other create index concurrently
- Takes longer to complete
- May fail
 - Results in invalid index
 - Try again



Fix index bloat: Summary

- There are some ways to fix index bloat
- In common:
 - Do not repair a bloated index
 - Create new / replace old
 - Require manual action
- Differences:
 - Locking
 - Level of convenience
- Is this really what we were looking for?



REINDEX, like autovacuum?

- We like to automate everything
- We like to have a tool that:
 - Monitors index bloat
 - Automatically fixes index bloat
- Should we build one, ourselves?
- How do you deal with index bloat?



Conclusions

- Index bloat can get pretty bad
 - Demo: Almost 8000%
- Ways to prevent index bloat
 - Tweak fillfactor / autovacuum
- Ways to fix index bloat
 - REINDEX, CREATE INDEX CONCURRENTLY
- Index bloat should be monitored



Thank you!